

OpenShift Virtualization VM Networking

Dale Bewley

Specialist SA

NA West OpenShift  Team

Red Hat



Multus Meta CNI Plugin

The CNI used by OpenShift

Is a meta-plugin which executes other plugins

Creates the interfaces in the pod used by the VM



Container Network Interface

CNI Plugins enable pod network configuration

```
{  
  "cniVersion": "0.3.1",  
  "name": "mynet",  
  "type": <plug>  
  <parameters>...  
}
```



Bridge



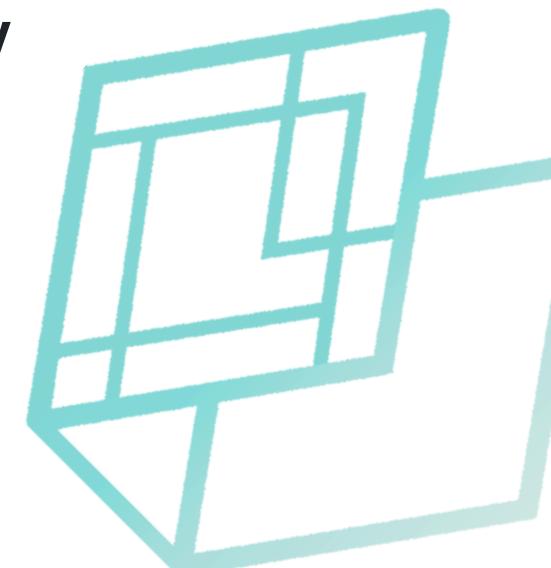
ovn-k8s-cni-overlay



SR-IOV



...





CNI Plugins

There are many plugins and they may be chained by Multus

```
$ oc rsh -n openshift-multus \
multus-2tt2c \
ls -1 /var/lib/cni/bin
bandwidth
bond
bridge
cert approver
cnv bridge
cnv tuning
dhcp
dummy
egress router
firewall
```

```
host-device
host-local
install_multus
ip-control-loop
ipvlan
kubeconfig_generator
loopback
macvlan
multus
multus-daemon
multus-shim
network-passt-binding
ovn-k8s-cni-overlay
...
```

 **bridge Plugin**

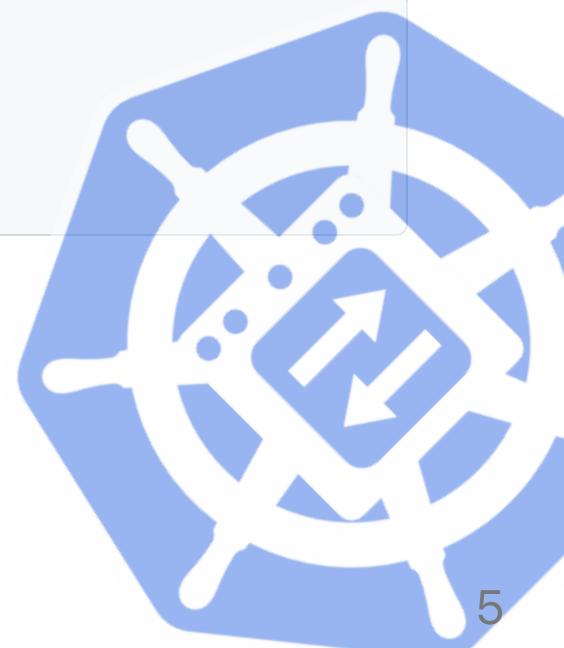
Don't use cnv-bridge, use bridge. They are identical.

Bridge can be used to attach to a Linux Bridge which may enable VLAN access or [VGT](#)

```
{  
  "cniVersion": "0.3.1",  
  "name": "mynet",  
  "type": "bridge",  
  ...  
}
```

🔌 ovn-k8s-cni-overlay Plugin

```
{  
  "cniVersion": "0.3.1",  
  "name": "mynet",  
  "type": "ovn-k8s-cni-overlay",  
  "topology": 🌐 <topology>,  
  ...  
}
```



Discussion of network attachments available for Virtual Machines on OpenShift with example implementations

KubeVirt Interfaces and Networks

 Primary Cluster Network

 Primary User Defined Networks

 Secondary Localnet (VLANS)

 Secondary User Defined Networks

VM Examples - Cluster Network

All VMs have the same address internally and masquerade using the pod IP.

```
# VM attached to default cluster network
apiVersion: kubevirt.io/v1
kind: VirtualMachine
metadata:
  name: vm-on-cluster-network
spec:
  template:
    spec:
      domain:
        devices:
          interfaces:
            - macAddress: '02:86:5e:00:00:07'
              masquerade: {}
              model: virtio
              name: default
      networks:
        - name: default
          pod: {}
```

Virt-Launcher Pod

Two ethernet interfaces in the virt launcher pod.

- Infrastructure locked `10.128.0.0/14` cluster network
- Always on `10.0.2.1/24`
- `k6t-eth0` short for `kubevirt-eth0` is a bridge enslaving `tap0`
- `tap0` is passed to QEMU for `eth0` in the VM

```
sh-5.1$ ip -c link
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: eth0@if379: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue state UP mode DEFAULT group default
    link/ether 0a:58:0a:83:01:61 brd ff:ff:ff:ff:ff:ff link-netnsid 0
3: k6t-eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue state UP mode DEFAULT group default qlen 1000
    link/ether 02:00:00:00:00:00 brd ff:ff:ff:ff:ff:ff
4: tap0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc fq_codel master k6t-eth0 state UP mode DEFAULT group default qlen 1000
    link/ether be:53:ae:c8:c5:66 brd ff:ff:ff:ff:ff:ff

sh-5.1$ ip -br -c -4 a
lo          UNKNOWN      127.0.0.1/8
eth0@if379   UP          10.131.1.97/23
k6t-eth0     UP          10.0.2.1/24
```

Virtual Machine

- VM *always* has IP `10.0.2.2/24`.
- Masquerades as pod IP `10.131.1.97/23` above.

```
[cloud-user@vm-pod ~]$ ip -c link
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc fq_codel state UP mode DEFAULT group default qlen 1000
    link/ether 02:86:5e:00:00:10 brd ff:ff:ff:ff:ff:ff
        altname enp1s0

[cloud-user@vm-pod ~]$ ip -br -c -4 a
lo          UNKNOWN      127.0.0.1/8
eth0         UP          10.0.2.2/24

[cloud-user@vm-pod ~]$ ip -c route
default via 10.0.2.1 dev eth0 proto dhcp src 10.0.2.2 metric 100
10.0.2.0/24 dev eth0 proto kernel scope link src 10.0.2.2 metric 100
```

VM on Default Cluster Network

Summary

Virt-launcher Pods

- eth0 on cluster network `10.128.0.0/14`
- k6t-eth0 always has IP `10.0.2.1/24`

VirtualMachines

- eth0 is always IP `10.0.2.2/24`
- Default gateway is always `10.0.2.1` on virt-launcher pod
- Masquerades at pod edge as IP of the virt-launcher pod
- Masquerades at node edge as IP of node default interface `br-ex`

VM Examples - Primary UDN

VMs have unique IPs from UDN subnet

Only Layer2 topology is supported
(localnet soon)

```
apiVersion: k8s.ovn.org/v1
kind: UserDefinedNetwork
  name: primary-udn
spec:
  topology: Layer2
  layer2:
    ipam:
      lifecycle: Persistent
    role: Primary
  subnets:
    - 10.1.1.0/24
```

...continued

```
# VM attached to primary UDN
apiVersion: kubevirt.io/v1
kind: VirtualMachine
metadata:
  name: vm-on-primary-udn
spec:
  template:
    spec:
      domain:
        devices:
          interfaces:
            - binding:
                name: l2bridge
                model: virtio
                name: default
      networks:
        - name: default
          pod: {}
```

Virt-Launcher Pod

Two ethernet interfaces in the virt launcher pod.

- Infrastructure locked `10.128.0.0/14` cluster network for kubelet health checks *only*
- Unique IP on the UDN range `10.1.1.0/24`

```
sh-5.1$ ip -c link
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: eth0@if356: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue state UP mode DEFAULT group default
    link/ether 0a:58:0a:83:01:4b brd ff:ff:ff:ff:ff:ff link-netnsid 0
3: ovn-udn1-nic@if357: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue master k6t-ovn-udn1 state UP mode DEFAULT group default
    link/ether 06:1b:c3:df:4d:d3 brd ff:ff:ff:ff:ff:ff link-netnsid 0
4: k6t-ovn-udn1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue state UP mode DEFAULT group default qlen 1000
    link/ether 06:1b:c3:df:4d:d3 brd ff:ff:ff:ff:ff:ff
5: tap0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc fq_codel master k6t-ovn-udn1 state UP mode DEFAULT group default qlen 1000
    link/ether 82:38:45:e8:1a:3d brd ff:ff:ff:ff:ff:ff
6: ovn-udn1: <BROADCAST,NOARP> mtu 1400 qdisc noop state DOWN mode DEFAULT group default qlen 1000
    link/ether 0a:58:0a:01:01:03 brd ff:ff:ff:ff:ff:ff

sh-5.1$ ip -br -c -4 a
lo          UNKNOWN      127.0.0.1/8
eth0@if356   UP          10.131.1.75/23
ovn-udn1     DOWN        10.1.1.3/24
```

Virtual Machine

One ethernet interface in the VM with IP from primary UDN

```
[cloud-user@vm-primary-udn ~]$ ip -c link
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc fq_codel state UP mode DEFAULT group default qlen 1000
    link/ether 0a:58:0a:01:01:03 brd ff:ff:ff:ff:ff:ff
        altname enp1s0

[cloud-user@vm-primary-udn ~]$ ip -br -c -4 a
lo          UNKNOWN      127.0.0.1/8
eth0         UP           10.1.1.3/24

[cloud-user@vm-primary-udn ~]$ ip -c route
default via 10.1.1.1 dev eth0 proto dhcp src 10.1.1.3 metric 100
10.1.1.0/24 dev eth0 proto kernel scope link src 10.1.1.3 metric 100
```

VM on Primary User Defined Network

Virt-Launcher Pod

- Two ethernet interfaces
- eth0@if356 is on infrastructure locked cluster network `10.128.0.0/14`
- ovn-udn1 is on primary UDN `10.1.1.3/24`

Virtual Machine

- eth0 has unique IP `10.1.1.3/24` from primary UDN of this Namespace
- Default gateway is `10.1.1.1`
- Masquerades at UDN gateway router as the IP from `169.254.0.0/17` associated with the UDN
- Masquerades at node edge as IP of node default interface `br-ex`

Masquerade Subnet

Each UDN has two IPs allocated from the masquerade subnet `169.254.0.0/17`.

```
# oc get network.operator/cluster -o yaml
apiVersion: operator.openshift.io/v1
kind: Network
metadata:
  name: cluster
spec:
  clusterNetwork:
  - cidr: 10.128.0.0/14
    hostPrefix: 23
  defaultNetwork:
    ovnKubernetesConfig:
      egressIPConfig: {}
      gatewayConfig:
        ipv4: {} # <-- default: 169.254.0.0/17
        ipv6: {} # <-- default: fd69::/112
        routingViaHost: false
```

VM Examples - Primary and Secondary UDN

```
apiVersion: k8s.ovn.org/v1
kind: UserDefinedNetwork
  name: primary-udn
spec:
  topology: Layer2
  layer2:
    ipam:
      lifecycle: Persistent
    role: Primary
  subnets:
    - 10.1.1.0/24
```

```
apiVersion: k8s.ovn.org/v1
kind: UserDefinedNetwork
  name: secondary-udn
spec:
  topology: Layer2
  layer2:
    ipam:
      lifecycle: Persistent
    role: Secondary
  subnets:
    - 10.2.2.0/24
```

...continued

```
# VM attached to primary UDN and secondary UDN
apiVersion: kubevirt.io/v1
kind: VirtualMachine
metadata:
  name: vm-on-primary-udn
spec:
  template:
    spec:
      domain:
        devices:
          interfaces:
            - binding:
                name: l2bridge
                model: virtio
                name: default
            - bridge: {}
                macAddress: '02:86:5e:00:00:0a'
                model: virtio
                name: secondary-udn
      networks:
        - name: default
          pod: {}
        - multus:
            networkName: secondary-udn
            name: secondary-udn
```

Virt-Launcher Pod

```
sh-5.1$ ip -c link
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
  link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: eth0@if412: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue state UP mode DEFAULT group default
  link/ether 0a:58:0a:83:01:80 brd ff:ff:ff:ff:ff:ff link-netnsid 0
3: ovn-udn1-nic@if413: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue master k6t-ovn-udn1 state UP mode DEFAULT group default
  link/ether 92:16:66:87:e3:d3 brd ff:ff:ff:ff:ff:ff link-netnsid 0
4: 2eae7330186-nic@if414: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue master k6t-2eae7330186 state UP mode DEFAULT group default
  link/ether 26:fc:0d:92:fe:71 brd ff:ff:ff:ff:ff:ff link-netnsid 0
5: k6t-ovn-udn1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue state UP mode DEFAULT group default qlen 1000
  link/ether 92:16:66:87:e3:d3 brd ff:ff:ff:ff:ff:ff
6: tap0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc fq_codel master k6t-ovn-udn1 state UP mode DEFAULT group default qlen 1000
  link/ether ea:a8:63:f5:7c:f7 brd ff:ff:ff:ff:ff:ff
7: ovn-udn1: <BROADCAST,NOARP> mtu 1400 qdisc noop state DOWN mode DEFAULT group default qlen 1000
  link/ether 0a:58:0a:01:01:03 brd ff:ff:ff:ff:ff:ff
8: k6t-2eae7330186: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue state UP mode DEFAULT group default qlen 1000
  link/ether 26:fc:0d:92:fe:71 brd ff:ff:ff:ff:ff:ff
9: tap2eae7330186: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc fq_codel master k6t-2eae7330186 state UP mode DEFAULT group default qlen 1000
  link/ether 32:64:42:71:34:80 brd ff:ff:ff:ff:ff:ff
10: pod2eae7330186: <BROADCAST,NOARP> mtu 1400 qdisc noop state DOWN mode DEFAULT group default qlen 1000
   link/ether 02:00:0a:02:02:03 brd ff:ff:ff:ff:ff:ff

sh-5.1$ ip -br -c -4 a
lo          UNKNOWN      127.0.0.1/8
eth0@if412    UP         10.131.1.128/23
ovn-udn1     DOWN        10.1.1.3/24
k6t-2eae7330186  UP        169.254.75.11/32
pod2eae7330186  DOWN       10.2.2.1/24
```

Virtual Machine

```
[cloud-user@vm-secondary-udn ~]$ ip -c link
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc fq_codel state UP mode DEFAULT group default qlen 1000
    link/ether 0a:58:0a:01:01:03 brd ff:ff:ff:ff:ff:ff
        altname enp1s0
3: eth1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc fq_codel state UP mode DEFAULT group default qlen 1000
    link/ether 02:00:0a:02:02:03 brd ff:ff:ff:ff:ff:ff
        altname enp2s0

[cloud-user@vm-secondary-udn ~]$ ip -br -c -4 a
lo          UNKNOWN      127.0.0.1/8
eth0         UP           10.1.1.3/24
eth1         UP           10.2.2.1/24

[cloud-user@vm-secondary-udn ~]$ ip -c route
default via 10.1.1.1 dev eth0 proto dhcp src 10.1.1.3 metric 100
10.1.1.0/24 dev eth0 proto kernel scope link src 10.1.1.3 metric 100
10.2.2.0/24 dev eth1 proto kernel scope link src 10.2.2.1 metric 101
```

Summary VM on Primary & Secondary User Defined Network

Virt-Launcher Pod

- Three ethernet interfaces
- eth0@if412 is on infrastructure locked cluster network `10.128.0.0/14`
- ovn-udn1 is on primary UDN `10.1.1.3/24`
- pod2eae7330186 is on secondary UDN `10.2.2.1/24`

Virtual Machine

- eth0 has unique IP `10.1.1.3/24` from primary UDN of this Namespace
- eth1 has unique IP `10.2.2.1/24` from secondary UDN of this Namespace
- Default gateway is `10.1.1.1`
- Masquerades at UDN gateway router as the IP from `169.254.0.0/17` for the UDN
- Masquerades at node edge as IP of node default interface `br-ex`

VM Examples - Localnet Secondary

VMs directly attached to VLAN

No UDN support for localnet topology as of 4.18.9.

Net-attach-def only.

```
apiVersion: k8s.cni.cncf.io/v1
kind: NetworkAttachmentDefinition
metadata:
  name: demo-vlan-1924
spec:
  config: |-  
    {  
      "cniVersion": "0.4.0",  
      "name": "demo-vlan-1924",  
      "type": "ovn-k8s-cni-overlay",  
      "topology": "localnet",  
      "netAttachDefName": "demo-vm-localnet/demo-vlan-1924",  
      "vlanID": 1924,  
      "ipam": {}  
    }
```

```
# VM attached to secondary localnet
apiVersion: kubevirt.io/v1
kind: VirtualMachine
metadata:
  name: vm-localnet
spec:
  template:
    spec:
      domain:
        devices:
          interfaces:
            - macAddress: '02:86:5e:00:00:13'
              masquerade: {}
              model: virtio
              name: default
            - bridge: {}
              macAddress: '02:86:5e:00:00:14'
              model: virtio
              name: nic-vlan-1924
      networks:
        - name: default
          pod: {}
        - multus:
            networkName: demo-vlan-1924
            name: nic-vlan-1924
```

Virt-Launcher Pod

```
sh-5.1$ ip -c link
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: eth0@if607: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue state UP mode DEFAULT group default
    link/ether 0a:58:0a:83:00:54 brd ff:ff:ff:ff:ff:ff link-netnsid 0
3: 16711a0a730-nic@if608: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue master k6t-16711a0a730 state UP mode DEFAULT group default
    link/ether 12:6f:ef:f4:36:9a brd ff:ff:ff:ff:ff:ff link-netnsid 0
4: k6t-eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue state UP mode DEFAULT group default qlen 1000
    link/ether 02:00:00:00:00:00 brd ff:ff:ff:ff:ff:ff
5: tap0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc fq_codel master k6t-eth0 state UP mode DEFAULT group default qlen 1000
    link/ether 8a:47:bc:d4:8b:d0 brd ff:ff:ff:ff:ff:ff
6: k6t-16711a0a730: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue state UP mode DEFAULT group default qlen 1000
    link/ether 12:6f:ef:f4:36:9a brd ff:ff:ff:ff:ff:ff
7: tap16711a0a730: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc fq_codel master k6t-16711a0a730 state UP mode DEFAULT group default qlen 1000
    link/ether 7a:f1:ab:eb:e9:e1 brd ff:ff:ff:ff:ff:ff
8: pod16711a0a730: <BROADCAST,NOARP> mtu 1400 qdisc noop state DOWN mode DEFAULT group default qlen 1000
    link/ether 02:86:5e:00:00:16 brd ff:ff:ff:ff:ff:ff

sh-5.1$ ip -br -c -4 a
lo          UNKNOWN      127.0.0.1/8
eth0@if607   UP          10.131.0.84/23
k6t-eth0     UP          10.0.2.1/24
```

Virtual Machine

```
[cloud-user@vm-localnet ~]$ ip -c link
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc fq_codel state UP mode DEFAULT group default qlen 1000
    link/ether 02:86:5e:00:00:15 brd ff:ff:ff:ff:ff:ff
        altname enp1s0
3: eth1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc fq_codel state UP mode DEFAULT group default qlen 1000
    link/ether 02:86:5e:00:00:16 brd ff:ff:ff:ff:ff:ff
        altname enp2s0

[cloud-user@vm-localnet ~]$ ip -br -c -4 a
lo          UNKNOWN      127.0.0.1/8
eth0         UP          10.0.2.2/24
eth1         UP          192.168.4.71/24

[cloud-user@vm-localnet ~]$ ip -c route
default via 10.0.2.1 dev eth0 proto dhcp src 10.0.2.2 metric 100
default via 192.168.4.1 dev eth1 proto dhcp src 192.168.4.71 metric 101
10.0.2.0/24 dev eth0 proto kernel scope link src 10.0.2.2 metric 100
192.168.4.0/24 dev eth1 proto kernel scope link src 192.168.4.71 metric 101
```

Summary VM on Localnet

Virt-launcher Pod

- eth0@if607 is on the cluster network `10.128.0.0/14`
- tap0 is passed to QEMU for eth0
- tap16711a0a730 is passed to QEMU for eth1

Virtual Machine

- eth0 is always IP `10.0.2.2/24`
- Default gateway is always `10.0.2.1` on virt-launcher
- Masquerades at node edge as IP of node default interface `br-ex`
- eth1 is `192.168.4.71/24` from DHCP on datacenter VLAN 1924

VM Examples - Localnet Secondary UDN

VMs directly attached to VLAN

TBD



Support for UDN localnet topology is targeted for 4.19.

Q&A / further reading

- [CNI Spec](#)

References

CNI spec & OVN-Kubernetes overview

2. OVN-Kubernetes (Primary CNI)

- **ovn-kubernetes GitHub**
 - Website: ovn-kubernetes.io/
 - Repo: [ovn-org/ovn-kubernetes](https://github.com/ovn-org/ovn-kubernetes)
 - Core repo for the OVN-Kubernetes integration: CNI binaries, controllers, docs.
- **OpenShift OVN-Kubernetes Guide**
 - Docs: [OpenShift Container Platform Networking – OVN-Kubernetes](#)
 - Red Hat's overview of