

Deep dive into CNI

Presentation Outline

Intro & objectives

CNI spec & OVN-Kubernetes overview

2. OVN-Kubernetes (Primary CNI)

- **ovn-kubernetes GitHub**
 - Website: ovn-kubernetes.io/
 - Repo: [ovn-org/ovn-kubernetes](https://github.com/ovn-org/ovn-kubernetes)
 - Core repo for the OVN-Kubernetes integration: CNI binaries, controllers, docs.
- **OpenShift OVN-Kubernetes Guide**
 - Docs: [OpenShift Container Platform Networking – OVN-Kubernetes](#)
 - Red Hat's overview of

Multus architecture

Bridge plugin deep dive (config + demo)

MAC VLAN plugin deep dive

Debugging tools & workflows

KubeVirt Interfaces and Networks

- Masquerade on Cluster Network
- L2 Bridge on Primary User Defined Network
-

VM Examples - Cluster Network

All VMs have the same address internally and masquerade using the pod IP.

```
# VM attached to default cluster network
apiVersion: kubevirt.io/v1
kind: VirtualMachine
metadata:
  name: vm-on-cluster-network
spec:
  template:
    spec:
      domain:
        devices:
          interfaces:
            - macAddress: '02:86:5e:00:00:07'
              masquerade: {}
              model: virtio
              name: default
      networks:
        - name: default
          pod: {}
```

Virt-Launcher Pod

Two ethernet interfaces in the virt launcher pod.

- Infrastructure locked `10.128.0.0/14` cluster network
- Always on `10.0.2.1/24`

```
sh-5.1$ ip -c link
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: eth0@if379: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue state UP mode DEFAULT group default
    link/ether 0a:58:0a:83:01:61 brd ff:ff:ff:ff:ff:ff link-netnsid 0
3: k6t-eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue state UP mode DEFAULT group default qlen 1000
    link/ether 02:00:00:00:00:00 brd ff:ff:ff:ff:ff:ff
4: tap0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc fq_codel master k6t-eth0 state UP mode DEFAULT group default qlen 1000
    link/ether be:53:ae:c8:c5:66 brd ff:ff:ff:ff:ff:ff

sh-5.1$ ip -br -c -4 a
lo                UNKNOWN      127.0.0.1/8
eth0@if379        UP            10.131.1.97/23
k6t-eth0          UP            10.0.2.1/24
```

Virtual Machine

Ethernet interface in the VM *always* has IP `10.0.2.2/24` . Masquerades as pod IP `10.131.1.97/23` above.

```
[cloud-user@vm-pod ~]$ ip -c link
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc fq_codel state UP mode DEFAULT group default qlen 1000
    link/ether 02:86:5e:00:00:10 brd ff:ff:ff:ff:ff:ff
    altname enp1s0

[cloud-user@vm-pod ~]$ ip -br -c -4 a
lo                UNKNOWN          127.0.0.1/8
eth0              UP                10.0.2.2/24

[cloud-user@vm-pod ~]$ ip -c route
default via 10.0.2.1 dev eth0 proto dhcp src 10.0.2.2 metric 100
10.0.2.0/24 dev eth0 proto kernel scope link src 10.0.2.2 metric 100
```

VM on Default Cluster Network

Summary

Virt-launcher Pods

- eth0 on cluster network `10.128.0.0/14`
- k6t-eth0 always has IP `10.0.2.1/24`

VirtualMachines

- eth0 is always IP `10.0.2.2/24`
- Default gateway is always `10.0.2.1` on virt-launcher pod
- Masquerades at pod edge as IP of the virt-launcher pod
- Masquerades at node edge as IP of node default interface `br-ex`

VM Examples - Primary UDN

VMs have unique IPs from UDN subnet

Layer2 topology only (`localnet` soon)

```
apiVersion: k8s.ovn.org/v1
kind: UserDefinedNetwork
  name: primary-udn
spec:
  topology: Layer2
  layer2:
    ipam:
      lifecycle: Persistent
      role: Primary
      subnets:
        - 10.1.1.0/24
```

```
# VM attached to primary UDN
apiVersion: kubevirt.io/v1
kind: VirtualMachine
metadata:
  name: vm-on-primary-udn
spec:
  template:
    spec:
      domain:
        devices:
          interfaces:
            - binding:
                name: l2bridge
                model: virtio
                name: default
      networks:
        - name: default
          pod: {}
```

Virt-Launcher Pod

Two ethernet interfaces in the virt launcher pod.

- Infrastructure locked `10.128.0.0/14` cluster network for kubelet health checks *only*
- Unique IP on the UDN range `10.1.1.0/24`

```
sh-5.1$ ip -c link
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: eth0@if356: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue state UP mode DEFAULT group default
    link/ether 0a:58:0a:83:01:4b brd ff:ff:ff:ff:ff:ff link-netnsid 0
3: ovn-udn1-nic@if357: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue master k6t-ovn-udn1 state UP mode DEFAULT group default
    link/ether 06:1b:c3:df:4d:d3 brd ff:ff:ff:ff:ff:ff link-netnsid 0
4: k6t-ovn-udn1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue state UP mode DEFAULT group default qlen 1000
    link/ether 06:1b:c3:df:4d:d3 brd ff:ff:ff:ff:ff:ff
5: tap0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc fq_codel master k6t-ovn-udn1 state UP mode DEFAULT group default qlen 1000
    link/ether 82:38:45:e8:1a:3d brd ff:ff:ff:ff:ff:ff
6: ovn-udn1: <BROADCAST,NOARP> mtu 1400 qdisc noop state DOWN mode DEFAULT group default qlen 1000
    link/ether 0a:58:0a:01:01:03 brd ff:ff:ff:ff:ff:ff

sh-5.1$ ip -br -c -4 a
lo                UNKNOWN      127.0.0.1/8
eth0@if356        UP            10.131.1.75/23
ovn-udn1          DOWN         10.1.1.3/24
```

Virtual Machine

One ethernet interface in the VM with IP from primary UDN

```
[cloud-user@vm-primary-udn ~]$ ip -c link
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc fq_codel state UP mode DEFAULT group default qlen 1000
    link/ether 0a:58:0a:01:01:03 brd ff:ff:ff:ff:ff:ff
    altname enp1s0

[cloud-user@vm-primary-udn ~]$ ip -br -c -4 a
lo                UNKNOWN          127.0.0.1/8
eth0              UP                10.1.1.3/24

[cloud-user@vm-primary-udn ~]$ ip -c route
default via 10.1.1.1 dev eth0 proto dhcp src 10.1.1.3 metric 100
10.1.1.0/24 dev eth0 proto kernel scope link src 10.1.1.3 metric 100
```

VM on Primary User Defined Network

Summary

Virt-launcher Pods

- Two ethernet interfaces
- eth0@if356 is on infrastructure locked cluster network `10.128.0.0/14`
- ovn-udn1 is on primary UDN `10.1.1.3/24`

VirtualMachines

- eth0 has unique IP `10.1.1.3/24` from primary UDN of this Namespace
- Default gateway is `10.1.1.1`
- Masquerades at UDN gateway router as the IP from `169.254.0.0/17` associated with the UDN

Masquerade Subnet

Each UDN has two IPs allocated from the masquerade subnet.

```
# oc get network.operator/cluster -o yaml
apiVersion: operator.openshift.io/v1
kind: Network
metadata:
  name: cluster
spec:
  clusterNetwork:
  - cidr: 10.128.0.0/14
    hostPrefix: 23
  defaultNetwork:
    ovnKubernetesConfig:
      egressIPConfig: {}
      gatewayConfig:
        ipv4: {} # <-- default: 169.254.0.0/17
        ipv6: {} # <-- default: fd69::/112
        routingViaHost: false
```

VM Examples - Primary and Secondary UDN

```
apiVersion: k8s.ovn.org/v1
kind: UserDefinedNetwork
  name: primary-udn
spec:
  topology: Layer2
  layer2:
    ipam:
      lifecycle: Persistent
      role: Primary
      subnets:
        - 10.1.1.0/24
```

```
apiVersion: k8s.ovn.org/v1
kind: UserDefinedNetwork
  name: secondary-udn
spec:
  topology: Layer2
  layer2:
    ipam:
      lifecycle: Persistent
      role: Secondary
      subnets:
        - 10.2.2.0/24
```

VM Examples - Primary and Secondary UDN

```
# VM attached to primary UDN and secondary UDN
apiVersion: kubevirt.io/v1
kind: VirtualMachine
metadata:
  name: vm-on-primary-udn
spec:
  template:
    spec:
      domain:
        devices:
          interfaces:
            - binding:
                name: l2bridge
                model: virtio
                name: default
            - bridge: {}
              macAddress: '02:86:5e:00:00:0a'
              model: virtio
              name: secondary-udn
      networks:
        - name: default
          pod: {}
        - multus:
            networkName: secondary-udn
            name: secondary-udn
```

Virt-Launcher Pod

```
sh-5.1$ ip -c link
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: eth0@if412: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue state UP mode DEFAULT group default
    link/ether 0a:58:0a:83:01:80 brd ff:ff:ff:ff:ff:ff link-netnsid 0
3: ovn-udn1-nic@if413: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue master k6t-ovn-udn1 state UP mode DEFAULT group default
    link/ether 92:16:66:87:e3:d3 brd ff:ff:ff:ff:ff:ff link-netnsid 0
4: 2eae7330186-nic@if414: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue master k6t-2eae7330186 state UP mode DEFAULT group default
    link/ether 26:fc:0d:92:fe:71 brd ff:ff:ff:ff:ff:ff link-netnsid 0
5: k6t-ovn-udn1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue state UP mode DEFAULT group default qlen 1000
    link/ether 92:16:66:87:e3:d3 brd ff:ff:ff:ff:ff:ff
6: tap0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc fq_codel master k6t-ovn-udn1 state UP mode DEFAULT group default qlen 1000
    link/ether ea:a8:63:f5:7c:f7 brd ff:ff:ff:ff:ff:ff
7: ovn-udn1: <BROADCAST,NOARP> mtu 1400 qdisc noop state DOWN mode DEFAULT group default qlen 1000
    link/ether 0a:58:0a:01:01:03 brd ff:ff:ff:ff:ff:ff
8: k6t-2eae7330186: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue state UP mode DEFAULT group default qlen 1000
    link/ether 26:fc:0d:92:fe:71 brd ff:ff:ff:ff:ff:ff
9: tap2eae7330186: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc fq_codel master k6t-2eae7330186 state UP mode DEFAULT group default qlen 1000
    link/ether 32:64:42:71:34:80 brd ff:ff:ff:ff:ff:ff
10: pod2eae7330186: <BROADCAST,NOARP> mtu 1400 qdisc noop state DOWN mode DEFAULT group default qlen 1000
    link/ether 02:00:0a:02:02:03 brd ff:ff:ff:ff:ff:ff

sh-5.1$ ip -br -c -4 a
lo                UNKNOWN      127.0.0.1/8
eth0@if412        UP           10.131.1.128/23
ovn-udn1          DOWN        10.1.1.3/24
k6t-2eae7330186   UP          169.254.75.11/32
pod2eae7330186    DOWN        10.2.2.1/24
```

Virtual Machine

```
[cloud-user@vm-secondary-udn ~]$ ip -c link
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc fq_codel state UP mode DEFAULT group default qlen 1000
    link/ether 0a:58:0a:01:01:03 brd ff:ff:ff:ff:ff:ff
    altnam enp1s0
3: eth1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc fq_codel state UP mode DEFAULT group default qlen 1000
    link/ether 02:00:0a:02:02:03 brd ff:ff:ff:ff:ff:ff
    altnam enp2s0

[cloud-user@vm-secondary-udn ~]$ ip -br -c -4 a
lo                UNKNOWN          127.0.0.1/8
eth0              UP                10.1.1.3/24
eth1              UP                10.2.2.1/24

[cloud-user@vm-secondary-udn ~]$ ip -c route
default via 10.1.1.1 dev eth0 proto dhcp src 10.1.1.3 metric 100
10.1.1.0/24 dev eth0 proto kernel scope link src 10.1.1.3 metric 100
10.2.2.0/24 dev eth1 proto kernel scope link src 10.2.2.1 metric 101
```

VM on Primary & Secondary User Defined Network

Summary

Virt-launcher Pods

- Three ethernet interfaces
- eth0@if412 is on infrastructure locked cluster network `10.128.0.0/14`
- ovn-udn1 is on primary UDN `10.1.1.3/24`
- pod2eae7330186 is on secondary UDN `10.2.2.1/24`

VirtualMachines

- eth0 has unique IP `10.1.1.3/24` from primary UDN of this Namespace
- eth1 has unique IP `10.2.2.1/24` from secondary UDN of this Namespace
- Default gateway is `10.1.1.1`

- Multus log:

- Pod annotations

```
oc get pods -n demo-udn-2 virt-launcher-fedora-black-elephant-21-msldg -o json \
  | jq -r '.metadata.annotations."k8s.ovn.org/pod-networks"' | jq -s
...
```


Q&A / further reading

- [CNI Spec](#)